

CPSC 420 Review Session


- ▶ Final Exam: Tue Apr 18, 2023 08:30am LSK 200.
One 2-sided page of notes

Today's Plan (Yikes)

- ▶ Convex Hull & its Algorithms
- ▶ Voronoi diagrams
- ▶ Linear Programming
- ▶ Network Flow
- ▶ Ford-Fulkerson algorithm
- ▶ Maximum matching in bipartite graphs
- ▶ Linear programming duality
- ▶ Compression
- ▶ Huffman Coding & Lempel-Ziv Compression
- ▶ P, NP, and NP-hardness
- ▶ Karp's 21 Problems
- ▶ Approximation algorithms
- ▶ Hardness of approximation
- ▶ Online Algorithms
- ▶ Cuckoo Hashing
- ▶ RSA cryptosystem
- ▶ Quantum Computing
- ▶ Zero-knowledge Proofs


Convex Hull

A set S is **convex** if for all $a, b \in S$ the segment \overline{ab} is in S .



Convex Hull


A set S is **convex** if for all $a, b \in S$ the segment \overline{ab} is in S .



The **convex hull** of a set P of points is the intersection of all convex sets that contain P , or equivalently the set of all convex combinations of P ($\sum \lambda_i p_i, \sum \lambda_i = 1$).


Convex Hull

A set S is **convex** if for all $a, b \in S$ the segment \overline{ab} is in S .




The **convex hull** of a set P of points is the intersection of all convex sets that contain P , or equivalently the set of all convex combinations of P ($\sum \lambda_i p_i, \sum \lambda_i = 1$).



A point $p \in P$ is on the boundary of $\text{CH}(P)$ iff there exists a line ℓ through p with all P on one side of ℓ .






Convex Hull Algorithms



Convex Hull Algorithms



Convex Hull Algorithms



Reductions (the fast kind)

As a reminder: the concept of reductions is not only for NP! Here's a tiny refresher.

Reductions (the fast kind)

As a reminder: the concept of reductions is not only for NP! Here's a tiny refresher.

I can use convex hull to sort using the mapping $i \mapsto (i, i^2)$ in time equal to $O(\text{Convex} + n)$, so convex hull can't be faster than comparison sorting, and I can use a decision tree to show that's $\Omega(\lg n)$.

Reductions (the fast kind)

As a reminder: the concept of reductions is not only for NP! Here's a tiny refresher.

I can use convex hull to sort using the mapping $i \mapsto (i, i^2)$ in time equal to $O(\text{Convex} + n)$, so convex hull can't be faster than comparison sorting, and I can use a decision tree to show that's $\Omega(\lg n)$.

(Reminder) in essence the tree has $n!$ leaves, which requires a depth of $\Omega(n \lg n)$.

Voronoi diagrams

Definition

A **Voronoi diagram** of a set of n sites (points) s_1, \dots, s_n is a set of regions R_1, \dots, R_n where R_i is the set of points x such that $d(x, s_i) \leq d(x, s_j)$ for all j .

Voronoi diagrams

Definition

A **Voronoi diagram** of a set of n sites (points) s_1, \dots, s_n is a set of regions R_1, \dots, R_n where R_i is the set of points x such that $d(x, s_i) \leq d(x, s_j)$ for all j .

A **Voronoi edge** is the border between two regions:
 $\{x \mid x \in R_i \text{ and } x \in R_j \text{ and } i \neq j\}$.

Voronoi diagrams

Definition

A **Voronoi diagram** of a set of n sites (points) s_1, \dots, s_n is a set of regions R_1, \dots, R_n where R_i is the set of points x such that $d(x, s_i) \leq d(x, s_j)$ for all j .

A **Voronoi edge** is the border between two regions:
 $\{x | x \in R_i \text{ and } x \in R_j \text{ and } i \neq j\}$.

A **Voronoi vertex** is the intersection of Voronoi edges:
 $\{x | x \text{ in more than 2 Vor. Regions}\}$.

Voronoi diagrams

Definition


A **Voronoi diagram** of a set of n sites (points) s_1, \dots, s_n is a set of regions R_1, \dots, R_n where R_i is the set of points x such that $d(x, s_i) \leq d(x, s_j)$ for all j .

A **Voronoi edge** is the border between two regions:
 $\{x | x \in R_i \text{ and } x \in R_j \text{ and } i \neq j\}$.

A **Voronoi vertex** is the intersection of Voronoi edges:
 $\{x | x \text{ in more than 2 Vor. Regions}\}$.

Problem Given $S = s_1, s_2, \dots, s_n$, find Voronoi vertices and edges

Algorithm from Computational Geometry by Preparata & Shamos



Exercises

1. Of Graham and Jarvis, which is faster on a circle of points?

Exercises

1. Of Graham and Jarvis, which is faster on a circle of points?
Graham
2. Which is faster on a line?

Exercises

1. Of Graham and Jarvis, which is faster on a circle of points?
Graham
2. Which is faster on a line? **Jarvis**
3. How can you tell a vertex is on the convex hull using a Voronoi diagram?

Exercises

1. Of Graham and Jarvis, which is faster on a circle of points?
Graham
2. Which is faster on a line? **Jarvis**
3. How can you tell a vertex is on the convex hull using a Voronoi diagram? **It has infinite area**
4. Say I offered you a $O(h \log(\log h)^4)$ algorithm for Convex Hull. Is that cool?

Exercises

1. Of Graham and Jarvis, which is faster on a circle of points?
Graham
2. Which is faster on a line? **Jarvis**
3. How can you tell a vertex is on the convex hull using a Voronoi diagram? **It has infinite area**
4. Say I offered you a $O(h \log(\log h)^4)$ algorithm for Convex Hull. Is that cool? **No**

Linear Programming

Definition

A **Linear Program** is a problem that can be arranged into the form

$$\begin{array}{ll}\text{maximize} & c^T x \\ \text{subject to} & Ax \leq b \\ & x \geq 0\end{array}$$

Linear Programming

Definition

A **Linear Program** is a problem that can be arranged into the form

$$\begin{array}{ll}\text{maximize} & c^T x \\ \text{subject to} & Ax \leq b \\ & x \geq 0\end{array}$$

A popular algorithm for solving LPs is Danzig's **Simplex**

Algorithm:

1. Start at a vertex v of the feasible set
2. While there is a neighbor v' of v with better objective value
3. $v = v'$

Network Flows

A **flow network** is a directed graph $G = (V, E)$ in which each edge $(u, v) \in E$ has a positive **capacity** $c(u, v)$ (non-edges have capacity 0).

G contains a **source** vertex s and a **sink** vertex t .

Network Flows

A **flow network** is a directed graph $G = (V, E)$ in which each edge $(u, v) \in E$ has a positive **capacity** $c(u, v)$ (non-edges have capacity 0).

G contains a **source** vertex s and a **sink** vertex t .

A **flow** is an assignment f of real numbers to edges of G :

1. **Non-negativity:** $0 \leq f_e$
2. **Capacity** $f_e \leq c_e$
3. **Conservation** $v \neq s, t : \sum f_{\delta^-(v)} = \sum f_{\delta^+(v)}$

Network Flows

A **flow network** is a directed graph $G = (V, E)$ in which each edge $(u, v) \in E$ has a positive **capacity** $c(u, v)$ (non-edges have capacity 0).

G contains a **source** vertex s and a **sink** vertex t .

A **flow** is an assignment f of real numbers to edges of G :

1. **Non-negativity:** $0 \leq f_e$
2. **Capacity** $f_e \leq c_e$
3. **Conservation** $v \neq s, t : \sum f_{\delta^-(v)} = \sum f_{\delta^+(v)}$

The **size** (or **value**) of a flow is: $\text{size}(f) = \sum_v f(s, v) - f(v, s)$

Network Flows

A **flow network** is a directed graph $G = (V, E)$ in which each edge $(u, v) \in E$ has a positive **capacity** $c(u, v)$ (non-edges have capacity 0).


G contains a **source** vertex s and a **sink** vertex t .

A **flow** is an assignment f of real numbers to edges of G :

1. **Non-negativity:** $0 \leq f_e$
2. **Capacity** $f_e \leq c_e$
3. **Conservation** $v \neq s, t : \sum f_{\delta^-(v)} = \sum f_{\delta^+(v)}$



The **size** (or **value**) of a flow is: $\text{size}(f) = \sum_v f(s, v) - f(v, s)$

Goal: Find flow with maximum size.




Max Flow via Path Augmentation [Ford & Fulkerson 1962]

1. Start with zero flow (a feasible solution)
2. Repeat until impossible
 - ▶ Choose an **augmenting path** from s to t
 - ▶ Increase flow on this path as much as possible



Max Flow via Path Augmentation [Ford & Fulkerson 1962]

1. Start with zero flow (a feasible solution)
2. Repeat until impossible
 - ▶ Choose an **augmenting path** from s to t
 - ▶ Increase flow on this path as much as possible



The **residual network** of flow network $G = (V, E)$ with flow f is $G^f = (V, E^f)$ where


$$E^f = \{(u, v) \mid f(u, v) < c(u, v) \text{ or } f(v, u) > 0\}$$

The **residual capacity** of an edge $(u, v) \in E^f$ is


$$c^f(u, v) = \begin{cases} c(u, v) - f(u, v) & \text{if } f(u, v) < c(u, v) \\ f(v, u) & \text{if } f(v, u) > 0 \end{cases}$$

An **augmenting path** in G is an $s \rightsquigarrow t$ path in G^f


Residual Network Example




Flow network



Flow



Residual network



Ford & Fulkerson, cont.

More Terminology

A **cut** is a partition (S, T) of V such that $s \in S$ and $t \in T$. (Cut separates s from t .)

The **capacity** of cut (S, T) is $c(S, T) = \sum_{u \in S, v \in T} c(u, v)$

The **flow** across cut (S, T) is $f(S, T) = \sum_{u \in S, v \in T} f(u, v) - f(v, u)$

Ford & Fulkerson, cont.

More Terminology

A **cut** is a partition (S, T) of V such that $s \in S$ and $t \in T$. (Cut separates s from t .)

The **capacity** of cut (S, T) is $c(S, T) = \sum_{u \in S, v \in T} c(u, v)$

The **flow** across cut (S, T) is $f(S, T) = \sum_{u \in S, v \in T} f(u, v) - f(v, u)$

The Gem: Max-Flow Min-Cut

$$|f^*| = c(S^*, T^*)$$

Runtime

$O(m|f^*|)$ (pseudo-polynomial). **Solution?**

Ford & Fulkerson, cont.

More Terminology

A **cut** is a partition (S, T) of V such that $s \in S$ and $t \in T$. (Cut separates s from t .)

The **capacity** of cut (S, T) is $c(S, T) = \sum_{u \in S, v \in T} c(u, v)$

The **flow** across cut (S, T) is $f(S, T) = \sum_{u \in S, v \in T} f(u, v) - f(v, u)$

The Gem: Max-Flow Min-Cut

$$|f^*| = c(S^*, T^*)$$

Runtime

$O(m|f^*|)$ (pseudo-polynomial). **Solution?** *Edmonds-Karp* is $O(mn^2)$ by using shortest path. There is even faster though!



Maximum Matching in Bipartite Graphs

A **matching** in a graph G is a subset M of its edges with no vertex the endpoint of more than one edge in M .

A **maximum matching** is a matching with the maximum number of edges.

A **maximal matching** is a matching to which another edge cannot be added to form a new matching.

A **bipartite graph** is a graph $G = (V, E)$ where V can be partitioned into A and B such that $\forall (u, v) \in E$, either $u \in A$ and $v \in B$ or $u \in B$ and $v \in A$.



Maximum Matching in Bipartite Graphs



A **matching** in a graph G is a subset M of its edges with no vertex the endpoint of more than one edge in M .

A **maximum matching** is a matching with the maximum number of edges.

A **maximal matching** is a matching to which another edge cannot be added to form a new matching.

A **bipartite graph** is a graph $G = (V, E)$ where V can be partitioned into A and B such that $\forall (u, v) \in E$, either $u \in A$ and $v \in B$ or $u \in B$ and $v \in A$.

Given bipartite graph $G = (V, E)$ with partitions A and B
find maximum matching in G .



Exercises

Projects with Dependencies

We have n projects, each with value p_i . $p_i > 0$ is a profit, and $p_i < 0$ is a cost. Additionally, we have an acyclic dependency graph G where the edge (i, j) means i depends on j .

Goal: find projects to do so as to maximize profit.

LP Duality

The dual program is constructed as follows:

- ▶ Variable \longleftrightarrow Constraint
- ▶ Maximum \longleftrightarrow Minimum

$$\begin{array}{ll}\text{maximize} & c^T x \\ \text{subject to} & Ax \leq b \\ & x \geq 0\end{array}$$

$$\begin{array}{ll}\text{minimize} & b^T y \\ \text{subject to} & A^T y \geq c \\ & y \geq 0\end{array}$$

LP Duality

The dual program is constructed as follows:

- ▶ Variable \longleftrightarrow Constraint
- ▶ Maximum \longleftrightarrow Minimum

$$\begin{array}{ll}\text{maximize} & c^T x \\ \text{subject to} & Ax \leq b \\ & x \geq 0\end{array}$$

$$\begin{array}{ll}\text{minimize} & b^T y \\ \text{subject to} & A^T y \geq c \\ & y \geq 0\end{array}$$

Weak Duality gives us says the dual solution is always an upper bound on the primal solution. **Strong Duality** says that they have the same optima if it exists.

Dynamic Programming

Dynamic Programming is an oft-overloaded term. The two key things you need to call what you're doing *Dynamic Programming* are

Dynamic Programming

Dynamic Programming is an oft-overloaded term. The two key things you need to call what you're doing *Dynamic Programming* are

1. The solution is a combination of sub-problem solutions.

Dynamic Programming

Dynamic Programming is an oft-overloaded term. The two key things you need to call what you're doing *Dynamic Programming* are

1. The solution is a combination of sub-problem solutions.
2. There is a nice number of total unique subproblems, so we solve them once

Dynamic Programming

Dynamic Programming is an oft-overloaded term. The two key things you need to call what you're doing *Dynamic Programming* are

1. The solution is a combination of sub-problem solutions.
2. There is a nice number of total unique subproblems, so we solve them once

Within this, there are 2 key ways to approach the solution:

- ▶ *Recursive* or *Top-Down* solutions solve the top level by calling & memoizing sub-problems.
- ▶ *Iterative* or *Bottom-Up* solutions start with the sub-problems with no dependencies, then build up until they reach the top level.

Exercises: 0-1 Knapsack

We have n items, each with a profit p_i and a weight w_i . Given a total profit P and total weight W , does there exist a set $S \subseteq [n]$ with $\sum p_S \geq P, \sum w_S \leq W$?

Information Theory [Shannon 1948]



The **information content** / **surprisal** contained in a message i that has probability p_i of being sent is

$$\log_2 \frac{1}{p_i}$$

Entropy is the average information content of a message X :

$$H(X) = \sum_i p_i \log_2 \frac{1}{p_i} = - \sum_i p_i \log_2 p_i$$

Information Theory [Shannon 1948]



The **information content** / **surprisal** contained in a message i that has probability p_i of being sent is

$$\log_2 \frac{1}{p_i}$$

Entropy is the average information content of a message X :

$$H(X) = \sum_i p_i \log_2 \frac{1}{p_i} = - \sum_i p_i \log_2 p_i$$

Source Coding Theorem m i.i.d. random variables each with entropy $H(X)$ can be compressed into more than $mH(X)$ bits with negligible risk of information loss **but** using less than $mH(X)$ bits results almost certainly in information loss. [Wikipedia-ish]

Huffman Coding and LZ78

Huffman Coding

Given a set of characters $a_1, a_2, \dots, a_\alpha$ and a probability p_i for each a_i ($\sum_i p_i = 1$), construct an encoding c_i for each character a_i so that the expected length of an encoded message is minimized, then it's just lookup.

Huffman Coding and LZ78

Huffman Coding

Given a set of characters $a_1, a_2, \dots, a_\alpha$ and a probability p_i for each a_i ($\sum_i p_i = 1$), construct an encoding c_i for each character a_i so that the expected length of an encoded message is minimized, then it's just lookup.

LZ78

Parse input into distinct phrases reading from left to right.

Each phrase is the shortest string not already a phrase. The 0th phrase is \emptyset .

Output ic for each phrase w , where c is the last character of w and i is the index of phrase u where $w = u \circ c$

Length is $c(n)(\lg c(n) + \lg \alpha)$ bits, as good as any finite state compressor.

Exercises

1. Imagine I promise you 2 compression algorithms: the first compresses to 70% on average, and at worst it compresses to 110%. The second compresses to 25% on average, and at worst it compresses to 97%. Why would you ask for the first?

Exercises

1. Imagine I promise you 2 compression algorithms: the first compresses to 70% on average, and at worst it compresses to 110%. The second compresses to 25% on average, and at worst it compresses to 97%. Why would you ask for the first?
The second is impossible, so I am either wrong or lying.
2. T/F: A low probability message has high information

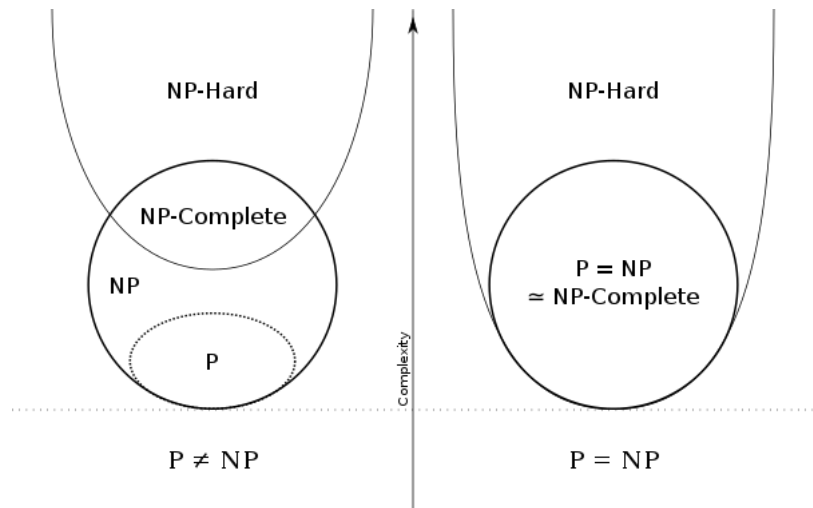
Exercises

1. Imagine I promise you 2 compression algorithms: the first compresses to 70% on average, and at worst it compresses to 110%. The second compresses to 25% on average, and at worst it compresses to 97%. Why would you ask for the first?
The second is impossible, so I am either wrong or lying.
2. T/F: A low probability message has high information **True**
3. Lempel-Ziv Compress the string AABAABBAA\$.

Exercises

1. Imagine I promise you 2 compression algorithms: the first compresses to 70% on average, and at worst it compresses to 110%. The second compresses to 25% on average, and at worst it compresses to 97%. Why would you ask for the first?
The second is impossible, so I am either wrong or lying.
2. T/F: A low probability message has high information **True**
3. Lempel-Ziv Compress the string AABAABBAA\$.
Dict:{,0A,1B,1A,0B,4A,1\$}, **String:** 0A1B1A0B4A1\$

Complexity Classes



What is a Reduction?

There are loads of ways to conceptualize a hardness reduction from $A \rightarrow B$:

- ▶ Prove that if I can solve B , I can solve A
- ▶ I can turn an input to A into an input for B , then transform the output back to the answer for A .
- ▶ If you hand me a B -machine, I can write a polytime algorithm for A

Reduction Structure & Strategy

Rules

The strategy of reduction we typically use is often called a *Karp Reduction* (guess why?). In this strategy, **you cannot mess with the oracle**. **All** you get to do is transform an input, then transform an output.

Proof Structure

With that in mind, a nice Completeness/Hardness answer for problem A usually follows the following structure (as do our rubrics):

Reduction Structure & Strategy

Rules

The strategy of reduction we typically use is often called a *Karp Reduction* (guess why?). In this strategy, **you cannot mess with the oracle**. **All** you get to do is transform an input, then transform an output.

Proof Structure

With that in mind, a nice Completeness/Hardness answer for problem A usually follows the following structure (as do our rubrics):

1. Prove A in NP **if** we want Completeness

Reduction Structure & Strategy

Rules

The strategy of reduction we typically use is often called a *Karp Reduction* (guess why?). In this strategy, **you cannot mess with the oracle**. **All** you get to do is transform an input, then transform an output.

Proof Structure

With that in mind, a nice Completeness/Hardness answer for problem A usually follows the following structure (as do our rubrics):

1. Prove A in NP **if** we want Completeness
2. State your input mapping and output mapping

Reduction Structure & Strategy

Rules

The strategy of reduction we typically use is often called a *Karp Reduction* (guess why?). In this strategy, **you cannot mess with the oracle**. **All** you get to do is transform an input, then transform an output.

Proof Structure

With that in mind, a nice Completeness/Hardness answer for problem A usually follows the following structure (as do our rubrics):

1. Prove A in NP **if** we want Completeness
2. State your input mapping and output mapping
3. Prove $A = T \implies B = T$ OR $B = F \implies A = F$

Reduction Structure & Strategy

Rules

The strategy of reduction we typically use is often called a *Karp Reduction* (guess why?). In this strategy, **you cannot mess with the oracle**. **All** you get to do is transform an input, then transform an output.

Proof Structure

With that in mind, a nice Completeness/Hardness answer for problem A usually follows the following structure (as do our rubrics):

1. Prove A in NP **if** we want Completeness
2. State your input mapping and output mapping
3. Prove $A = T \implies B = T$ OR $B = F \implies A = F$
4. Prove $A = F \implies B = F$ OR $B = T \implies A = T$

If it seems really intuitive **and** you think it's still rigorous enough, you can combine the last two, but beware!

Karp's 21 Problems

A quick bit of history

: In 1971 The Cook-Levin Theorem proves that SAT is NP-Hard from scratch, then in 1972 Karp uses reductions to prove the same for 21 other problems (a year after Edmonds-Karp). They receive a Turing Award each, and get everyone excited.

Karp's 21 Problems (Exercises)

Clique / Independent Set

- ▶ Set packing
- ▶ Vertex cover
 - ▶ Set covering
 - ▶ Feedback node set
 - ▶ Feedback arc set
 - ▶ Directed Hamilton cycle
 - ▶ Undirected Hamilton cycle

Oh, and 0–1 integer programming
(A sat-only variation)

3-SAT

- ▶ Chromatic number / Coloring
 - ▶ Clique cover
 - ▶ Exact cover
 - ▶ Hitting set
 - ▶ Steiner tree
 - ▶ 3-D matching
 - ▶ Knapsack / Subset Sum
 - ▶ - Job sequencing
 - ▶ - Partition
 - ▶ - Max Cut

Approximation Algorithms

Approximation Algorithms are useful for **optimization problems**, not decision problems.

An $f(n)$ -approximation algorithm A for a minimization problem has

$$\frac{A}{OPT} \leq f(n)$$

An $f(n)$ -approximation algorithm A for a maximization problem has

$$\frac{OPT}{A} \leq f(n)$$

Sometimes, you get problems that are even hard to approximate!

Online Algorithms

Online Algorithms work on streams of input instead of fixed inputs.
An $f(n)$ -competitive algorithm A for a minimization problem has

$$\frac{\mathbb{E}(A)}{OPT} \leq f(n)$$

An $f(n)$ -competitive algorithm A for a maximization problem has

$$\frac{\mathbb{E}(A)}{OPT} \geq f(n)$$

OPT here is **not** an online algorithm. It gets the whole string up front!

Wait, those equations look different.

Online Algorithms

Online Algorithms work on streams of input instead of fixed inputs.
An $f(n)$ -competitive algorithm A for a minimization problem has

$$\frac{\mathbb{E}(A)}{OPT} \leq f(n)$$

An $f(n)$ -competitive algorithm A for a maximization problem has

$$\frac{\mathbb{E}(A)}{OPT} \geq f(n)$$

OPT here is **not** an online algorithm. It gets the whole string up front!

Wait, those equations look different. How does this make me feel?

Online Algorithms

Online Algorithms work on streams of input instead of fixed inputs.
An $f(n)$ -competitive algorithm A for a minimization problem has

$$\frac{\mathbb{E}(A)}{OPT} \leq f(n)$$

An $f(n)$ -competitive algorithm A for a maximization problem has

$$\frac{\mathbb{E}(A)}{OPT} \geq f(n)$$

OPT here is **not** an online algorithm. It gets the whole string up front!

Wait, those equations look different. How does this make me feel? Confused.

Exercise

Knapsack 2-Approximation

We have n items, each with a profit p_i and a weight w_i . Given a total weight W , maximize $\sum p_S \geq P$ by picking $S \subseteq [n]$ with $\sum w_S \leq W$.

Randomized Marking Mouse

If Mouse follows a deterministic strategy, there is a sequence S of Cat probes that causes

$$\text{MouseCost}(S) \geq (m - 1)\text{OPT}(S)$$

Paging

$m - 1$ = cache size

m = different pages

Mouse = page not in cache

Cat probes = page requests

Must move = page fault

Randomized Marking Mouse (RMM)

- Start at random spot
- If Cat probes a spot, mark it
- If Cat probes Mouse's spot,
Mouse moves to random unmarked spot
- If Mouse is at last unmarked spot, clear marks [phase ends]

Randomized Marking Mouse performance

Claim: $E[\text{RMMCost}(S)] \leq O(\log m)\text{OPT}(S)$

Proof: Initially, RMM is equally likely to be at any of the m spots.

1st probe finds Mouse with probability $1/m$.

Whether Mouse is found or not, Mouse is at each of the $m-1$ unmarked spots with prob. $1/(m-1)$.

2nd probe (to unmarked spot) finds Mouse with prob $1/(m-1)$.

Mouse is at each of the $m-2$ unmarked spots with prob. $1/(m-2)$. Etc.

Let $X_i = \begin{cases} 1 & \text{if Mouse found on } i\text{th probe to unmarked spot} \\ 0 & \text{otherwise} \end{cases}$

$$\begin{aligned} E[\text{\#times found per phase}] &= E[X_1 + X_2 + \cdots + X_m] \\ &\leq \frac{1}{m} + \frac{1}{m-1} + \cdots + \frac{1}{1} = O(\log m) \end{aligned}$$

OPT moves once per phase.



Is Totally Random Mouse (TRM) better?

TRM runs to a random spot if found.

Consider the Methodical Cat (MC):

- Probe spots 1, 2, 3, ... until Mouse found
- Repeat

What does the OPT mouse do?

Is Totally Random Mouse (TRM) better?

TRM runs to a random spot if found.

Consider the Methodical Cat (MC):

- Probe spots 1, 2, 3, ... until Mouse found
- Repeat

What does the OPT mouse do? Hide in spot m

$$E[\text{\#times RM found before MC probes } m] = \\ E[\text{\#rolls of } m\text{-sided dice before } m] = m$$

\Rightarrow RM is m -competitive.

Random Marking Mouse is best

Claim: Any Mouse A has $E[A(S)] \in \Omega(\log m)\text{OPT}(S)$

Proof:

Idea: Show that a Cat exists that will cause $E[A(S)] \in \Omega(\log m)$ regardless of the Mouse.

Random Cat (RC) probes a random spot with each probe. RC finds Mouse with prob. $\frac{1}{m}$ no matter what Mouse does.
 $\Rightarrow E[A(S)]$ after t probes is $\frac{t}{m}$.

How many RC probes until RC examines every spot?

Coupon Collector Problem $\Rightarrow \Theta(m \log m)$

So OPT Mouse (that knows RC's probes) moves once in sequence S of $\Omega(m \log m)$ probes, while Mouse A moves $E[A(S)] \in \frac{\Omega(m \log m)}{m} = \Omega(\log m)$ times. □

Universal Families of Hash Functions

A family of hash functions H (that map $U \rightarrow \{0, 1, \dots, m-1\}$) is **universal** if for all distinct keys $x, y \in U$

$$\Pr_{h \in H}[h(x) = h(y)] \leq \frac{1}{m}.$$

Example

Let $h_{a,b}(x) = ((ax + b) \bmod p) \bmod m$ where p is a prime bigger than any key.

$$H = \{h_{a,b} | a \in \{1, 2, \dots, p-1\}, b \in \{0, 1, \dots, p-1\}\}$$

Cuckoo Hashing

Time per operation

Find $O(1)$ time worst case

Delete $O(1)$ time worst case

Insert $O(1)$ expected, amortized time

- ▶ Use two hash functions h_1 and h_2 .
- ▶ Item x will be stored in slot $h_1(x)$ or $h_2(x)$ of hash table T .
- ▶ Each slot in the hash table can contain at most one item.
- ▶ n = maximum number of items stored at any time
- ▶ m = size of hash table T ($m > n$)

On an **insert**(x) collision, item x kicks the resident item y out. Item y then goes to its alternate slot (kicking whoever's there out). Etc. Etc.

RSA public/private key cryptosystem [Rivest,Shamir,Adleman '77]

Bob has two functions: secret $S_B()$ and public $P_B()$

Properties:

1. $S_B(P_B(M)) = M$ and $P_B(S_B(M)) = M$
2. Hard to find M given $P_B(M)$ without $S_B()$

Alice sends $P_B(M)$ to Bob.

Bob decrypts: $S_B(P_B(M)) = M$

Good: Use again and again

Bad: No one knows if it's secure.

factoring easy \Rightarrow RSA breakable.

factoring hard \Rightarrow RSA secure? (unknown)

Digital Signatures: Alice sends $(M, \sigma = S_A(M))$ to Bob

Bob can check that $P_A(\sigma) = M$.

Quantum Computing

Qubits

A *qubit* is a combination $|0\rangle$ and $|1\rangle$:

$$|\psi\rangle = \alpha|0\rangle + \beta|1\rangle$$

α and β are both complex numbers here, but since it must be that $\alpha^2 + \beta^2 = 1$, we actually only have 3 degrees of freedom.

Gates

Quantum gates, at the end of the day, are unitary matrices. For example, Hadamard is $\frac{1}{\sqrt{2}} \begin{pmatrix} 1 & 1 \\ 1 & -1 \end{pmatrix}$.

Zero Knowledge

The key idea here, is prove your identity without giving away any of your secrets.

If there's time, a quick example: Colour-Blindness [Credits to Wikipedia]

Closing Remarks

Any Questions?