

CPSC 420 Lecture 23: Today's announcements:

- ▶ Examlet 3 on Mar 17 in class. Closed book & no notes
- ▶ Reading: Randomized Algorithms [by Motwani and Raghavan]

Today's Plan

- ▶ Online Algorithms
 - ▶ Hiring problem ✓
 - ▶ Page replacement
 - ▶ List Update
 - ▶ Experts

Online Algorithms

For input sequence p_1, p_2, \dots, p_n an **online algorithm** must produce an output given p_1, p_2, \dots, p_i for each i **without seeing** p_{i+1}, \dots, p_n .

Page replacement in a cache

p_1, \dots, p_n is a sequence of **page requests** made by a program.

k is cache size (in pages).

At the i th page request, p_i , the cache contains some k pages. If p_i is **not** in cache (**page fault**) some page must be evicted from cache to make room for p_i , then p_i is added to the cache.

The cost of a page replacement algorithm A on sequence p_1, p_2, \dots, p_n is

$$f_A(p_1, p_2, \dots, p_n) = \# \text{ faults } A \text{ has on } p_1, p_2, \dots, p_n$$

Online algorithm must decide what page to evict without knowing future requests.

Some paging algorithms

- LRU** (Least Recently Used) Evict page whose most recent request occurred furthest in the past.
- FIFO** (First In First Out) Evict page that has been in cache the longest.
- LFU** (Least Frequently Used) Evict page that has been requested least often.
- OPT** (Optimal) Evict best^{*} page knowing the future requests.

Some paging algorithms

- LRU** (Least Recently Used) Evict page whose most recent request occurred furthest in the past.
- FIFO** (First In First Out) Evict page that has been in cache the longest.
- LFU** (Least Frequently Used) Evict page that has been requested least often.
- OPT** (Optimal) Evict best^{*} page knowing the future requests.

* **next** request is furthest in the future

Paging algorithms: Example

page requests

time

	A	B	C	D	B	A	B	C	D	A	B	requests
--	---	---	---	---	---	---	---	---	---	---	---	----------

LRU

A	B	C	D	D	D	D	C	C	C	B
Y	A	B	C	C	A	A	A	D	D	D
Z	Y	A	B	B	B	B	B	B	A	A
*	*	*	*		*		*	*	*	*

9 faults

FIFO

A	B	C	D	D	A	B	C	D	A	B
Y	A	B	C	C	D	A	B	C	D	A
Z	Y	A	B	B	C	D	A	B	C	D
*	*	*	*		*	*	*	*	*	*

10 faults

LFU

A	B	C	D	D	D	D	C	C	A	A
Y	A	B	C	C	A	A	A	D	D	D
Z	Y	A	B	B	B	B	B	B	B	B
*	*	*	*		*		*	*	*	

8 faults

OPT

A	B	C	D	D	D	D	D	D	D	B
Y	A	B	B	B	B	B	C	C	C	C
Z	Y	A	A	A	A	A	A	A	A	A
*	*	*	*				*			*

6 faults

Evaluating online algorithms

Which algorithm is best?

Worst case cost

$$n = \max_{p_1 \dots p_n} f_{\text{LRU}}(p_1 \dots p_n) = \max_{p_1 \dots p_n} f_{\text{FIFO}}(p_1 \dots p_n) = \max_{p_1 \dots p_n} f_{\text{LFU}}(p_1 \dots p_n)$$

Every online algorithm can be made to fault on every request.

Average case cost

$$E[f_X(p_1 \dots p_n)] = n(1 - k/m)$$

where $m = \#$ possible pages

Competitive analysis

How does online alg compare to best offline alg?

An online algorithm A is c -**competitive** if for all p_1, p_2, \dots, p_n ,

$$f_A(p_1 \dots p_n) \leq c \cdot f_{\text{OPT}}(p_1 \dots p_n) + b$$

for some constant b .

Lower Bound on Competitive Factor

Claim: If A is a deterministic online alg for paging then $c \geq k$.

Proof: Idea: Find a sequence that is bad for A but good for OPT .

Suppose both A and OPT start with pages $1, 2, \dots, k$ in cache.

Request page $a_1 = k + 1$

A evicts some page, call it a_2

Request page a_2

A evicts some page, call it a_3

Request page a_3

etc...

\Rightarrow Online alg A faults on every request.

How many different pages are there? $k+1$

\nwarrow
 $k+1$ one
additional
page

Lower Bound on Competitive Factor

Claim: If A is a deterministic online alg for paging then $c \geq k$.

Proof: Idea: Find a sequence that is bad for A but good for OPT .
Suppose both A and OPT start with pages $1, 2, \dots, k$ in cache.

Request page $a_1 = k + 1$

A evicts some page, call it a_2

Request page a_2

A evicts some page, call it a_3

Request page a_3

etc...

\Rightarrow Online alg A faults on every request.

How many different pages are there? $k + 1$

Lower Bound on Competitive Factor

How often does OPT fault?

$$\frac{a_1 \ a_2 \ a_3 \ \dots \ a_j \ a_{j+1} \ \dots \text{requests}}{\text{OPT} \quad \quad \quad * \quad \quad \quad * \quad \quad \quad \text{faults}}$$

Sub-claim: OPT faults at most once in k successive page requests from this sequence

Proof.

OPT evicts the page p in cache that is requested furthest in the future. Since there are only $k + 1$ different pages, the next k pages requested can be kept in cache. □

$$\Rightarrow f_A(a_1, a_2, \dots, a_n) \geq k \cdot f_{\text{OPT}}(a_1, a_2, \dots, a_n) \quad \square$$

LRU is k -competitive

Theorem

LRU is k -competitive

Proof.

Let p_1, p_2, \dots, p_n be **any** sequence of page requests. Partition this sequence into contiguous subsequences (**phases**) such that LRU faults on the first page of the phase and the phase contains exactly k different pages (or $\leq k$ in the last phase).

$OPT \text{ faults} \geq 1$

$k = 3$	A C D C C B C A D A A requests					
LRU	*	?	?	*	*	*
OPT	?	?	?	*		*
	phase 1			ph 2		ph 3

$K \text{ #phases}$
 $\#phases - 1$

LRU faults $\leq k$ times per phase.

OPT must have the first page of a phase in cache at the beginning of a phase. Since the remainder of the phase plus the first page of the next phase consists of k different pages, OPT must fault at least once during these requests. $\Rightarrow OPT \text{ faults} \geq \#phases - 1$

Any Marking Algorithm is k -competitive

Marking Algorithm MARK

0. Start with all k pages in cache unmarked
1. On page request p
2. if p not in cache then
3. evict any unmarked page
(if no unmarked page, first unmark all k pages)
5. bring p into cache
6. mark p

	A	B	C	D	B	A	B	C	D	A	B
→ MARK	•A Y Z *	•A •B Z *	•A •B •C *	•D B C *	•D •B C *	•D •B •A *	•D •B •A *	•C B A *	•C •D A *	•C •D •A <div></div>	•B D A *
→ LRU	A Y Z *	B A Y *	C B A *	D C B *	D C B *	D A B *	D A B *	C A B *	C D B *	C D A *	B D A *

Any Marking Algorithm is k -competitive

Marking Algorithm MARK

0. Start with all k pages in cache unmarked
1. On page request p
2. if p not in cache then
3. evict any unmarked page
 (if no unmarked page, first unmark all k pages)
5. bring p into cache
6. mark p

Proof.

Partition p_1, p_2, \dots, p_n into **phases**, a maximal subsequence with k distinct pages. (The first starts with p_1 .) Assume p_1 is not in cache. MARK faults $\leq k$ times per phase.

OPT must have the first page p_i of a phase in cache at the beginning of a phase. Since the remainder of the phase plus the first page of the next phase consists of k different pages (different from p_i), OPT must fault at least once during these requests.

\Rightarrow OPT faults $\geq \# \text{phases} - 1$ times.



Randomized online algorithm

Online Hide and Seek


Mouse hides in one of m hiding places.

Cat looks in one hiding place each time step.

If Cat finds Mouse, Mouse runs to another place.

Cost = #times Mouse moves

OPT = min #times future-knowing Mouse must move



Randomized online algorithm

Online Hide and Seek


Mouse hides in one of m hiding places.

Cat looks in one hiding place each time step.

If Cat finds Mouse, Mouse runs to another place.

Cost = #times Mouse moves

OPT = min #times future-knowing Mouse must move



$$\text{OPT}(1\ 2\ 3\ 4\ 1\ 2\ 3\ 4) = 2$$

Randomized online algorithm

Online Hide and Seek


Mouse hides in one of m hiding places.

Cat looks in one hiding place each time step.

If Cat finds Mouse, Mouse runs to another place.

Cost = #times Mouse moves

OPT = min #times future-knowing Mouse must move



$$\text{OPT}(1\ 2\ 3\ 4\ 1\ 2\ 3\ 4) = 2$$

Randomized online algorithm

Online Hide and Seek


Mouse hides in one of m hiding places.

Cat looks in one hiding place each time step.

If Cat finds Mouse, Mouse runs to another place.

Cost = #times Mouse moves

OPT = min #times future-knowing Mouse must move



$$\text{OPT}(1\ 2\ 3\ 4\ 1\ 2\ 3\ 4) = 2$$

Randomized online algorithm

Online Hide and Seek


Mouse hides in one of m hiding places.

Cat looks in one hiding place each time step.

If Cat finds Mouse, Mouse runs to another place.

Cost = #times Mouse moves

OPT = min #times future-knowing Mouse must move



$$\text{OPT}(1\ 2\ 3\ 4\ 1\ 2\ 3\ 4) = 2$$

Randomized online algorithm

Online Hide and Seek


Mouse hides in one of m hiding places.

Cat looks in one hiding place each time step.

If Cat finds Mouse, Mouse runs to another place.

Cost = #times Mouse moves

OPT = min #times future-knowing Mouse must move



$$\text{OPT}(1\ 2\ 3\ 4\ 1\ 2\ 3\ 4) = 2$$

Randomized online algorithm

Online Hide and Seek


Mouse hides in one of m hiding places.

Cat looks in one hiding place each time step.

If Cat finds Mouse, Mouse runs to another place.

Cost = #times Mouse moves

OPT = min #times future-knowing Mouse must move



$$\text{OPT}(1\ 2\ 3\ 4\ 1\ 2\ 3\ 4) = 2$$

$$\text{OPT}(1\ 2\ 1\ 1\ 3\ 4\ 1\ 2) = 1$$